

You're almost there — sign up to start building in Notion today.

Sign up or login

Tata Sons HR — AI-First Hiring Platform

Tata Sons HR — AI-First Hiring Platform

A platform where **AI does the work**, humans approve the decisions, and everything is **auditable**.


A hiring platform that runs **end-to-end on AI** — from "we need to hire someone" to "the candidate has accepted" — with humans stepping in only at decision moments that need their judgment.

Three things make it different from anything off the shelf:

1. **The chat IS the product.** Hiring managers talk to an AI assistant in Microsoft Teams. They don't fill forms. They don't log into a portal. They just ask, in plain English, and the AI does the rest.
2. **Every AI decision is explainable.** Every score, every flag, every recommendation comes with a confidence number and a written reason. Humans can override anything in one click — and that override is permanent record.
3. **It's modular.** Eleven specialised AI workers, each doing one thing exceptionally well. We can upgrade one without touching the others. We can swap the AI provider without rewriting the platform.

How a manager actually uses it







The manager spent **about 15 minutes total** across the entire process. The platform spent the rest.

The architecture — three layers, clean separation

Each layer has **one job**, owns it completely, and is replaceable without touching the others. That's the whole game.








Who owns what — and what's swappable


Layer	What it owns	What can be swapped <i>(without code changes)</i>
1 · Experience	The conversation, the approval cards, the cockpit views	The channel — Teams today, WhatsApp / Slack / web tomorrow
2 · Platform	Workflow rules, audit trail, identity, confidence gating	Nothing. This is the durable IP. The compounding asset.
3 · AI Workers	Each AI capability — prompts, schemas, models	The model , prompt , even provider , per worker, independently

The key insight: the only "thick" boundary is between Layer 2 and Layer 3 — one HTTP envelope (`POST /aiworker/run`), eleven workers behind it. Add a twelfth worker tomorrow? One config change. Replace GPT-4 with Claude under the JD Writer? Commotion swaps it; our code never knows.

The full hiring workflow (and where each AI worker hooks in)








Legend: 🤖 blue = AI worker · 👤 amber = human decision · purple = workflow event

The manager doesn't think about workers. They just talk to the AI assistant — the right specialist gets called automatically.

How Commotion fits in

Commotion is our AI worker gateway. Our platform never talks to a model directly — every AI call goes through this contract.




Two separate quality nets — even if a model hallucinates or a worker drifts, malformed output never touches our database.

How Teams connects (and why this design)

Hiring managers live in Teams. The platform is **inside** Teams — not a tool that sends a Teams notification.






Why this design (not just "a notification bot"):

- **Approvals happen in chat, not via email links.** A JD approval card has Approve / Request Changes buttons — one tap, decision recorded, audit logged.
- **The AI assistant is conversational, not transactional.** Manager can revise mid-conversation, ask for changes in plain English, change their mind.
- **Identity flows through Microsoft.** Same single-sign-on as everything else. No new logins to remember.
- **Calendar invites are real Microsoft calendar events.** Interview slots become accepted invites with one tap.

The decision moment — when AI hands off to humans

Every AI decision has a confidence score (0–100%). Below a configurable threshold, the platform automatically pauses the workflow and asks the human in chat.



When confidence is low, a **review card** lands in the manager's Teams chat:

Red flag detected: Senior Backend Engineer • Asha S. | | | AI thinks: Concern (62% confidence) | | "3 stints under 1.5 years in last 4 years..." | | | [Approve] [Override (with reason)] |


One click — human in the loop, decision logged, audit-ready.

Thresholds are configurable per deployment, set by HR leadership, not engineers. Want red flags to go to humans more often? Move the threshold from 80% to 90%. Done in seconds, no deploy.

Why "modular" actually matters

A traditional hiring platform is a single monolith. Upgrading one capability means upgrading everything. Slow, risky, expensive.

Our architecture treats every AI capability as a **plug-in**:




Three concrete scenarios

Scenario	Without modularity	With our architecture
"Switch to a cheaper LLM"	Rewrite half the platform · weeks of work	Commotion swaps the model under one worker ID. Code unchanged. Minutes.
"Inclusivity audit found a missed pattern"	Prompt change touches every AI codepath · full QA cycle	Update one worker. The other 10 keep running. A morning.
"Legal asks for 7-year AI decision retention"	Find every AI touchpoint · change them all	Audit trail records every call through one chokepoint. Change one retention policy. An afternoon.

The four guardrails

What makes this safe to run on real hiring decisions, not just impressive in a demo:





Where we are today

Built	Tested	Documented
✓ All 11 AI worker contracts	✓ 275 automated tests	✓ Full prompt bundle for Commotion registration
✓ End-to-end platform (Teams bot, web cockpit, audit)	✓ Real call to Commotion's gateway returning structured output today	✓ Architecture decision log (17 decisions, public)
✓ Confidence-gated handover protocol	✓ Auth + envelope confirmed working	✓ Testing playbook (local + real Microsoft accounts)
✓ Bias guards in every worker prompt	✓ Error mapping (5 distinct failure modes)	✓ Worker registration spec for Commotion onboarding


What we need from Tata Sons — to swap our placeholders for real content


The platform is **built and tested end-to-end**, but a handful of pieces are running on our **best-guess placeholders** today. These are the items where Tata Sons' real institutional knowledge has to come in — they cannot be invented by engineers.

We've engineered every one of these as a **drop-in swap**: nothing below requires code changes, only configuration / content updates. Most can be filled in by the right HR, Brand, or IT contact in a single working session.

The asks, grouped by who owns them







What we need	Our placeholder today	Why it matters
--------------	-----------------------	----------------